

ResumeRoast

Template Library v1 — MVP Draft for Review

How It Works

Stage 1 already extracts **roast_targets[]** — 2-3 verbatim fragments from the resume (e.g. "Title: Aspiring Software Developer", "Currently Learning: Django", "NoteSphere: deployed, no user count stated").

Stage 2's new job: **pick the best-matching template, fill the slots with those fragments. That's it.** No creative generation. The humor comes from the template structure; the specificity comes from their actual resume text.

Model parameters: temp 0.7 (slot-filling, not creativity), max_tokens 100.

Specificity Demo (Sneha's actual targets)

Targets: ["Title: Aspiring Software Developer", "Currently Learning: Django", "NoteSphere: deployed, no user count stated"]

Template	Output
T04 — The POV	<i>POV: you shipped NoteSphere, listed Django as currently learning, and applied for a Software Developer role anyway. This is that resume.</i>
T05 — The Slow Load	<i>NoteSphere shipped. Django still buffering. The completion percentage is not on the resume.</i>

CATEGORY 1 — Currently Learning

TRIGGER: Fires when targets contain "Currently Learning: X"

T01 — The Backlog

{COURSE} is on the resume. {PROJECT} shipped. The loading indicator is rotating in the wrong direction.

{COURSE} = Currently Learning item | {PROJECT} = project name

T02 — The Inventory

Skills: {SKILL_COUNT} listed. Currently Learning: {LEARN_COUNT} more. One of these sections is growing faster.

{SKILL_COUNT} = count from skills section | {LEARN_COUNT} = count of learning items

T03 — The Timestamp

Title: {TITLE}. Currently Learning: {COURSE}. {COURSE} has been in production somewhere since {YEAR}.

{TITLE} = job title | {COURSE} = tech being learned | {YEAR} = hardcoded: Docker 2013, React 2013, Python 1991, Git 2005

T04 — The POV

POV: you shipped {PROJECT}, listed {COURSE} as currently learning, and applied for a {TITLE} role anyway. This is that resume.

{PROJECT} = project name | {COURSE} = learning item | {TITLE} = job title

T05 — The Slow Load

{PROJECT} shipped. {COURSE} still buffering. The completion percentage is not on the resume.

{PROJECT} = project name | {COURSE} = Currently Learning item

CATEGORY 2 — No Metrics

TRIGGER: Fires when targets contain "deployed, no user count" or "zero metrics" or "no numbers"

T06 — The Ledger

{PROJECT} shipped. Users: unconfirmed. Revenue: undisclosed. Impact: left as an exercise for the recruiter.

{PROJECT} = project name

T07 — The Filed Under

Built {PROJECT}. Documented the metrics under private information.

{PROJECT} = project name

T08 — The Receipt

{COUNT} projects shipped. Zero receipts attached. This is a portfolio the way a list of restaurants is a food review.

{COUNT} = number of projects on resume (Stage 1 can count these)

T09 — The Confidence

Deployed {PROJECT} with no usage numbers attached. The confidence required to submit this is its own qualification.

{PROJECT} = project name

T10 — The Portfolio

{PROJECT} exists. Whether anyone opened it is not addressed. Currently learning {COURSE} to maybe find out.

{PROJECT} = project name | {COURSE} = Currently Learning item

CATEGORY 3 — Aspiring Title

TRIGGER: Fires when targets contain "Aspiring" or "Enthusiast" in the title

T11 — The Job Description

Title: {TITLE}. {PROJECT} deployed. Aspirations and deliverables are listed on the same document with no apparent awareness of the irony.

{TITLE} = full title ("Aspiring Software Developer") | {PROJECT} = project name

T12 — The Honest Version

The honest version of this resume leads with {PROJECT}. The submitted version leads with {TITLE}.

{PROJECT} = best project | {TITLE} = aspirational title

T13 — The Forecast

{TITLE} is the current status. {PROJECT} is the evidence. The gap between those two things is the whole resume.

{TITLE} = aspirational title | {PROJECT} = project name

CATEGORY 4 — Self-Contradicting Project

TRIGGER: Fires when project name implies a utility the person hasn't demonstrated (job tracker with gaps, etc.)

T14 — The App That Tracks

Built {PROJECT} to solve a problem. The problem is still visible in the rest of the resume.

{PROJECT} = project name (best when it implies productivity/organization)

T15 — The Founder

Founded {COMPANY}. Revenue: unpublished. Clients: unnamed. {COMPANY} appears on the resume the way a band appears on someone's bio.

{COMPANY} = company/agency name from resume

T16 — The Stack

Full-stack project. Front end deployed. Back end: listed in skills. Evidence of the connection: not submitted.

No slots — fires when resume lists full-stack skills but project descriptions are thin

CATEGORY 5 — Skills Overload

TRIGGER: Fires when skills section has many items with little evidence of use

T17 — The Collection

Skills: {SKILL_A}, {SKILL_B}, {SKILL_C}, and {COUNT} others. Project that uses all of them: not found.

{SKILL_A/B/C} = first 3 skills listed | {COUNT} = total skills minus 3

T18 — The Bucket List

Skills section contains {COUNT} items. Currently Learning section contains {COUNT_2} more. At some point this is a bucket list.

{COUNT} = skills count | {COUNT_2} = learning items count

T19 — The Proficiency

Listed {SKILL} as proficient. {SKILL} appears nowhere in the projects or experience. It is proficient in theory.

{SKILL} = a skill claimed but absent from project descriptions

CATEGORY 6 — High-Signal Catchall

TRIGGER: Fires when any strong target pair exists — fallback category when others don't fit cleanly

T20 — The Juxtaposition

{FACT_1}. {FACT_2}. One of these is impressive. The other one is on the same page.

{FACT_1} = strongest positive target | {FACT_2} = most embarrassing target

T21 — The Diagnosis

{PROJECT} exists. Metrics do not. This is not a portfolio entry, it is an alibi.

{PROJECT} = project name with no numbers

T22 — The Filing System

Everything that required doing is on this resume. Everything that happened after doing it is not.

No slots — fires when resume has many projects but zero outcomes documented

T23 — The Progression

First project: {PROJECT_1}. Second project: {PROJECT_2}. Learning velocity: confirmed.

Documentation velocity: zero.

{PROJECT_1} and {PROJECT_2} = two project names from resume

T24 — The Status Update

{PROJECT}: shipped. {COURSE}: in progress. {TITLE}: current status. This resume is a loading screen.

{PROJECT} = project | {COURSE} = learning item | {TITLE} = job title

T25 — The Literal One

{TARGET_1}. {TARGET_2}. This is the whole argument for and against this candidate, on the same page.

{TARGET_1} and {TARGET_2} = the two starkest targets — one decent, one not

New Stage 2 Prompt Structure

The new Stage 2 system prompt will be dramatically shorter. Core structure:

"You are a slot-filler. You receive a template and resume targets. Fill the slots with the EXACT text from the targets. Output only the filled template — no quotes, no explanation, no additions."

The user message will include:

1. A subset of the best-matching templates (picked by a matching function, not the model)
2. The verbatim roast_targets from Stage 1
3. Instruction: pick one, fill slots, output it.

Temperature drops from 0.9 to **0.7**. Max tokens stays at **100**. The matching logic (which template category fires) is deterministic code, not AI judgment.

Pending Before Code Changes

PRIORITY	ITEM
CRITICAL	Push all current changes to GitHub — nothing committed yet, one crash = everything lost
APPROVE	Template list — remove any you hate, flag any missing categories
DECIDE	T03 year-lookup: hardcode a table (Docker 2013, React 2013, etc.) or skip this template?
DECIDE	T02 / T17 / T18 need counts — Stage 1 must extract skill_count and learning_count, or we skip these
AFTER	Task #3: fix verdict field (normalize.ts line 476 — still AI-controlled)
AFTER	Task #4: fix fake auth (auth/page.tsx handleSubmit just does router.push)
AFTER	Delete lib/normalize_clean.ts (confirmed dead file)